



Eidgenössische Technische Hochschule Zürich
Swiss Federal Institute of Technology Zurich

Computer Networks

Assignment 1: Reliable Transport

Reliable Transport

Implement a **reliable packet stream** (not byte stream!)

... which handles:

- Packet drops
- Packet corruption
- Flow control
- Packet reordering

Fundamental Mechanisms

Error detection

- Corrupt packets must be discarded
- Implemented via checksum

Acknowledgments (ACK)

- Small control packet to confirm the reception of a packet
- When a sender gets an ACK, sender learns that recipient has successfully gotten all packets until (excl.) ackno

Fundamental Mechanisms

Time-outs

- If the sender doesn't get an ACK after “reasonable” time, it retransmits the original packet

Naive Approach: Stop-and-Wait

Algorithm

- After transmitting one packet, sender waits for an ACK
- If the ACK doesn't arrive in time, sender retransmits

Disadvantage

- Inefficient use of link's capacity

Sliding Window Protocol

Objective: better utilization of link bandwidth by sending multiple unacknowledged packets



Send window (SW): set of unacknowledged packets

Receiver window (RW): set of received packets

Requirement: need to keep SW and RW synchronized

Sliding Window Definitions

“Frame” - This is a packet in the assignment, though typically it is bytes




Each frame is indivisible and has a sequence number (seqno)

Sliding Window: Sender



Sliding Window: Sender



**Sent and
acknowledged**

Sliding Window: Sender

Sent not yet
acknowledged



Sent and
acknowledged

Sliding Window: Sender

Sent not yet
acknowledged



Sent and
acknowledged

Not sent

SND.UNA: lowest seqno of outstanding frames



SND.UNA: lowest seqno of outstanding frames

The acknowledgement number (ackno)
is the next seqno the receiver expects to get



↑
SND.UNA

Acknowledgements are cumulative

$$\text{SND.UNA} = \max(\text{SND.UNA}, \text{ackno})$$

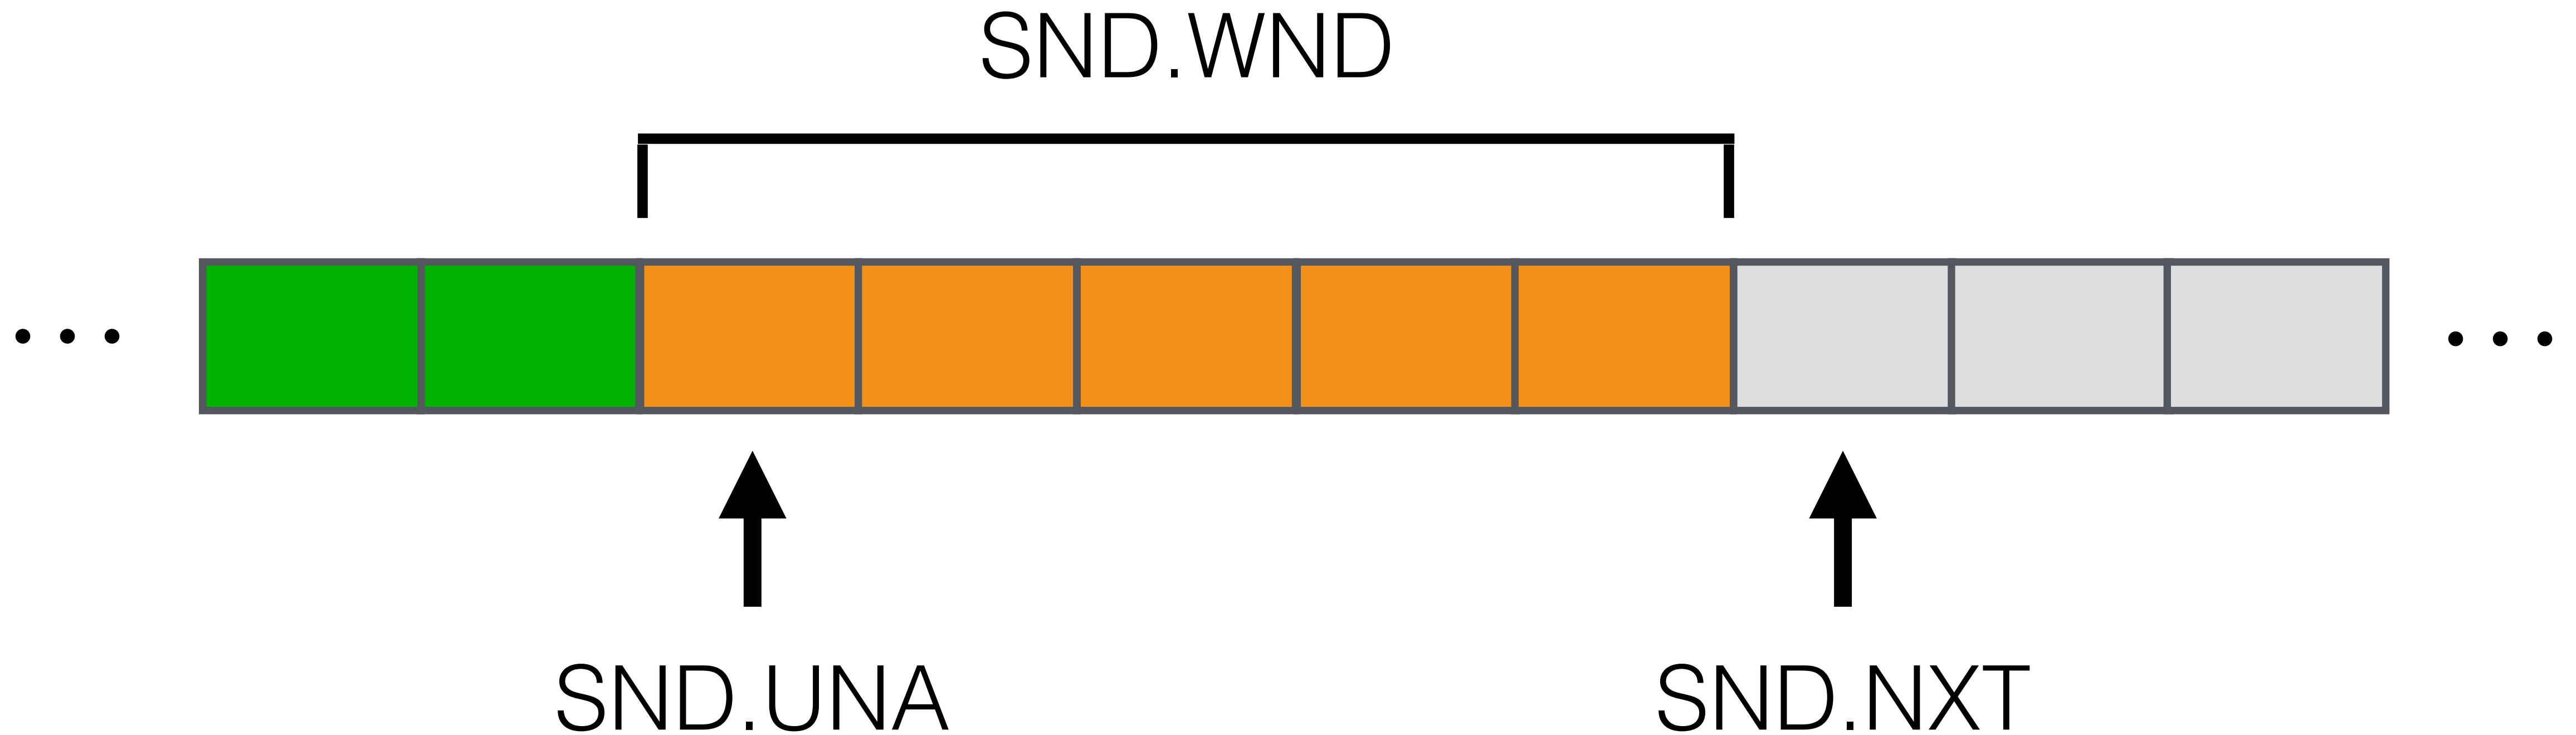
when an ACK arrives

**We don't do selective acknowledgement
in this assignment**

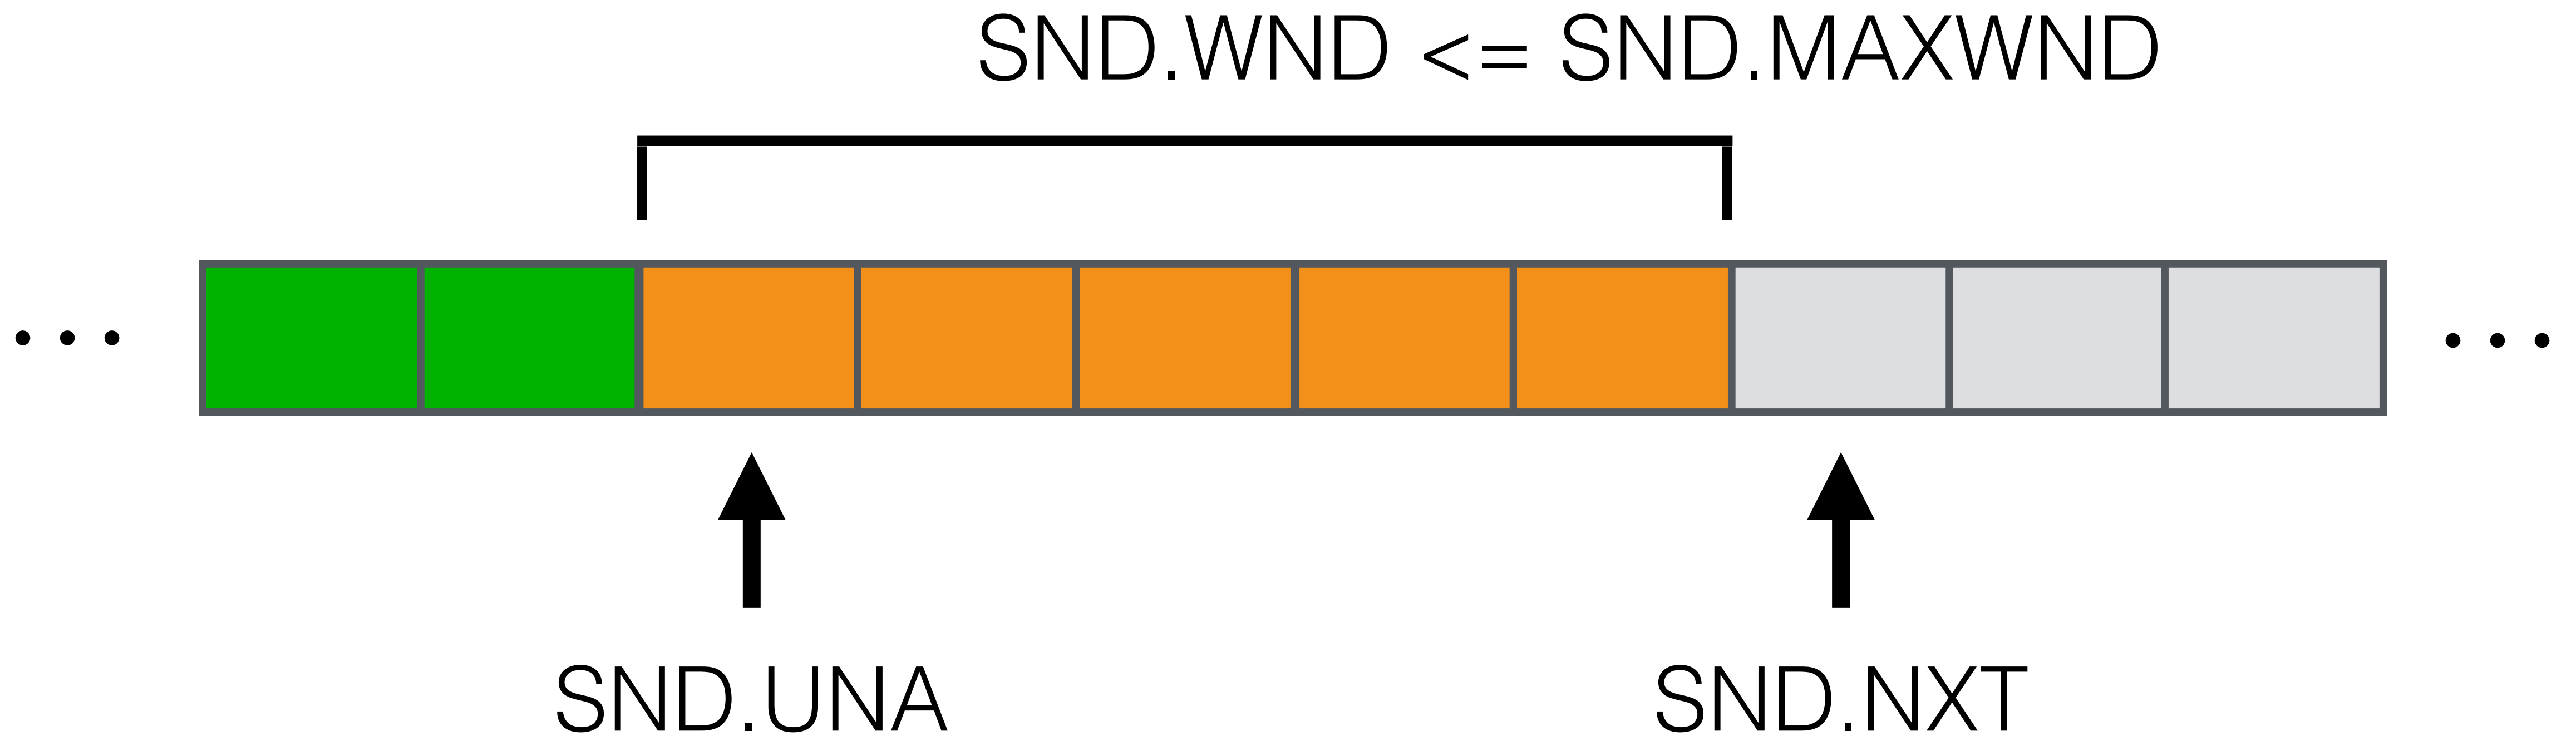
SND.NXT: seqno of next frame to send out




$$\text{SND.WND} = \text{SND.NXT} - \text{SND.UNA}$$



$SND.WND \leq SND.MAXWND$



Timeouts



Associate timeouts with each frame sent
Retransmit if no ACK received before timeout

Sender summary


- **State:** SND.UNA, SND.NXT
- $\text{SND.WND} = \text{SND.NXT} - \text{SND.UNA}$
- Send window is not a constant!
- $\text{SND.WND} \leq \text{SND.MAXWND}$
- Upon reception of ACK(ackno): $\text{SND.UNA} = \max(\text{SND.UNA}, \text{ackno})$
- Timeouts resend frames within the SND.WND
- You can only send a next “new” frame with $\text{seqno} = \text{SND.NXT}$ if $\text{SND.NXT} < \text{SND.UNA} + \text{SND.MAXWND}$
- After sending frame with $\text{seqno} = \text{SND.NXT}$, $\text{SND.NXT} = \text{SND.NXT} + 1$

Sliding Window: Receiver




RCV.NXT: next seqno expected

Received everything
till here




RCV.NXT: next seqno expected

Out-of-order frames received



RCV.WND: max. # of out-of-order frames it will accept



Receiver summary

- **State:** RCV.NXT
- RCV.WND = RCV.MAXWND

Receiver summary: upon receiving frame

If $\text{seqno} \geq \text{RCV.NXT} + \text{RCV.WND}$

Drop frame

Else

Store in the buffer if not already there

If $\text{seqno} == \text{RCV.NXT}$:

Set RCV.NXT to the highest seqno consecutively stored in the buffer + 1

Release data $[\text{seqno}, \text{RCV.NXT} - 1]$ to application layer

Send back ACK with cumulative $\text{ackno} = \text{RCV.NXT}$